



Week16_Course

Database System Concurrency Control part4_1

Shared locks Protocol

Database System - Nankai



Concurrency Control



- Schedule
- Conflict Serializable Schedule
- Two phase locking
- Warning Protocol
- Validation

Database System - Nankai



- Beyond this simple 2PL protocol, it is all a matter of improving performance and allowing more concurrency...
 - Shared locks
 - Multiple granularity
 - Inserts, deletes and phantoms
 - Other types of C.C. mechanisms
 - Timestamping
 - Validation



Shared locks

So far only exclusive locks:

$S = \dots l_1(A) r_1(A) u_1(A) \dots l_2(A) r_2(A) u_2(A) \dots$

Do not conflict \rightarrow locking unnecessary

Instead, use **shared** locks (S) for reading:

$S = \dots ls_1(A) r_1(A) ls_2(A) r_2(A) \dots u_1(A) u_2(A)$



Write actions conflict → use **exclusive (X)** locks for writing

Lock actions:

$l-m_k(A)$: lock A in mode m (S or X) for T_k

$u_k(A)$: release (whatever) lock(s) held by transaction T_k on element A



Rule #1 Well-formed transactions

$T_i = \dots l - S_1(A) \dots r_1(A) \dots u_1(A) \dots$

$T_i = \dots l - X_1(A) \dots w_1(A) \dots u_1(A) \dots$

- Request
 - an S-lock for reading
 - an X-lock for writing
- Release the locks eventually



- What about transactions that first read and later write the same element?

Option 1: Request exclusive lock

$T_i = \dots l-X_1(A) \dots r_1(A) \dots | w_1(A) \dots u_1(A) \dots$

Option 2: Upgrade

(E.g., need to read, but don't know if will write...)

$T_i = \dots l-S_1(A) \dots r_1(A) \dots l-X_1(A) \dots w_1(A) \dots u(A) \dots$

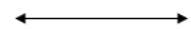
$T_j = [l-S_2(A)] \dots r_2(A)$

Think s getting
2nd lock on A



Rule #2 Legal scheduler

$S = \dots l-S_i(A) \dots \dots u_i(A) \dots$



$\text{no } l-X_j(A) \text{ for } j \neq i$

$S = \dots l-X_i(A) \dots \dots u_i(A) \dots$



$\text{no } l-X_j(A) \text{ for } j \neq i$

$\text{no } l-S_j(A) \text{ for } j \neq i$



A way to summarize Rule #2

Compatibility matrix:

		Lock requested by some T_j		
		S	X	C
Locks already held by some T_i	S	true	false	x
	X	false	x	x
	C	x	x	x

(A)

T_j

True \Leftrightarrow OK to give a new lock of requested kind



Rule # 3 (2PL)

Only change to previous:
Lock upgrades

$S(A) \rightarrow \{S(A), X(A)\}$ or

$S(A) \rightarrow X(A)$

are allowed only in the growing phase



Theorem Rules 1,2,3 \Rightarrow Conf.serializable
for S/X locks schedules

Proof: Similar to the X locks case

Lock types beyond S/X

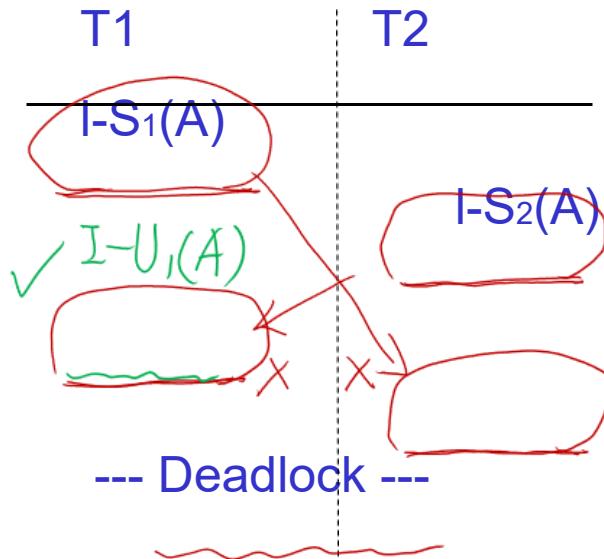
Examples:

- (1) update lock
- (2) increment lock (see the textbook)



Update locks

A common deadlock problem with upgrades:



Option 2: Upgrade

(E.g., need to read, but don't know if will write...)

T_i=... **I-S₁(A)** ... **r₁(A)** ... **I-X₁(A)** ... **w₁(A)** ... **u₁(A)** ...

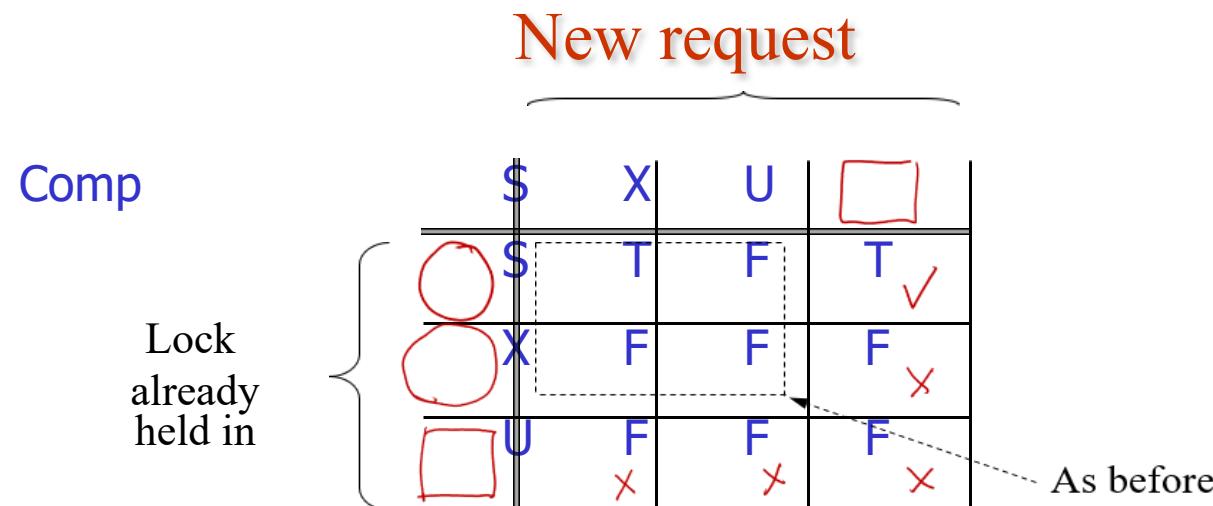
U
A

Database System - Nankai



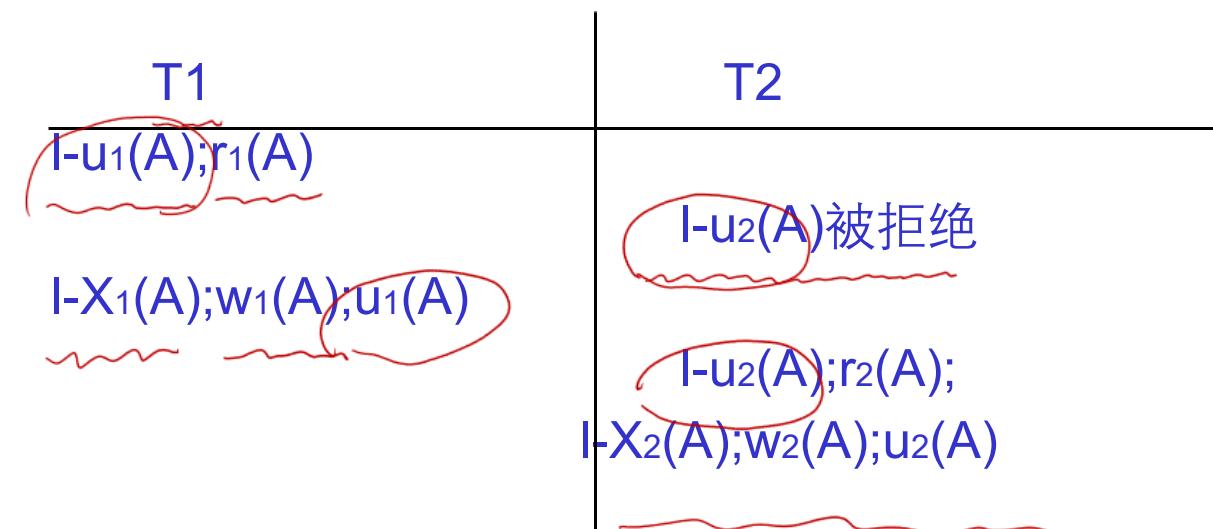
Solution

If T_i wants to read A and knows it may later want to write A, it requests an update lock (not shared) $l-u_k(A)$





A common deadlock problem with update:





Note: object A may be locked in different modes at the same time...

$S_1 = \dots I - S_1(A) \dots I - S_2(A) \dots I - U_3(A) \dots$

- To grant a lock in mode m, mode m must be compatible with all currently held locks on object



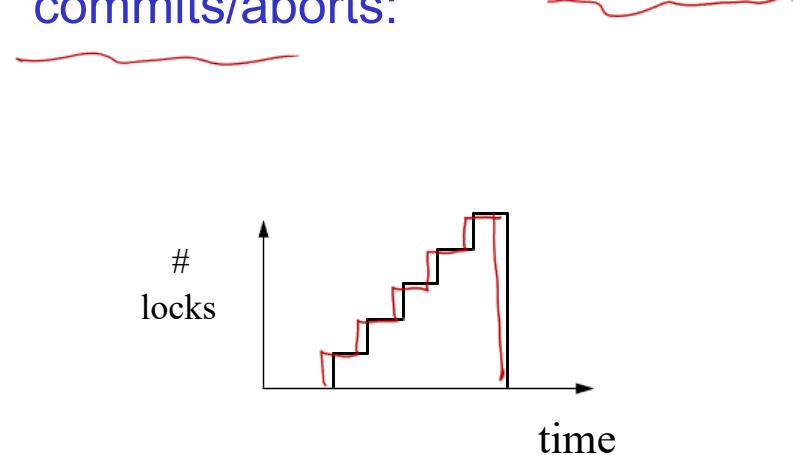
How does locking work in practice?

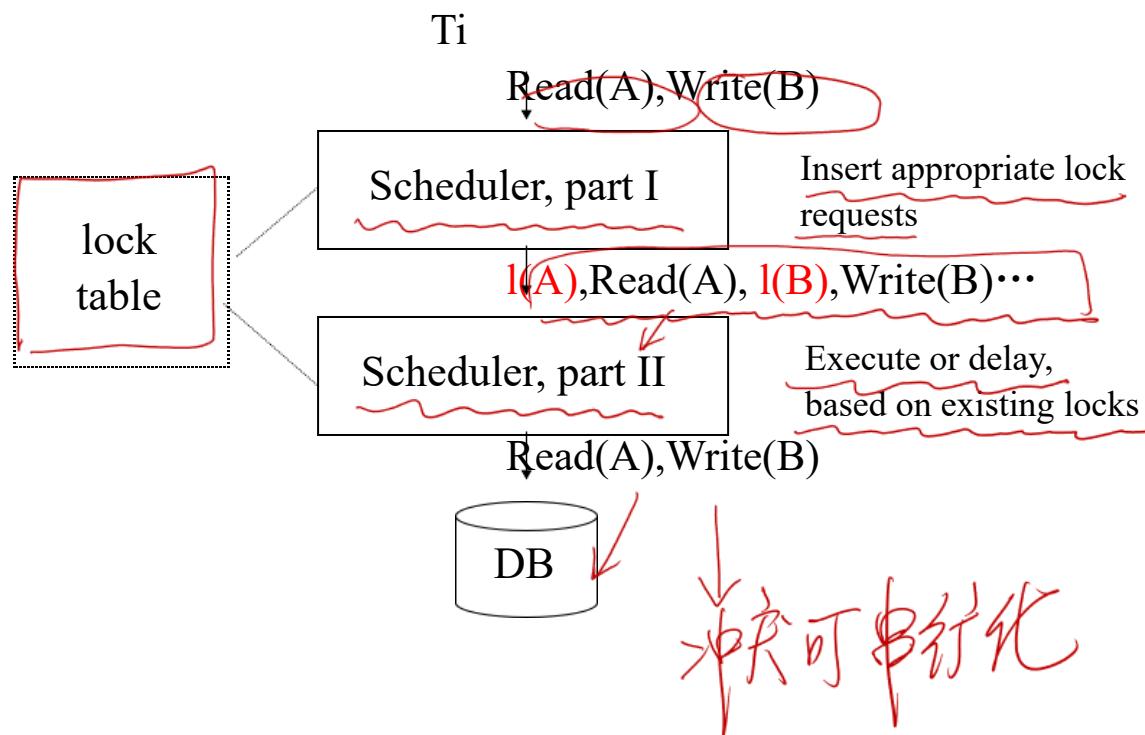
- Every system is different
(E.g., may not even provide
CONFLICT-SERIALIZABLE schedules)
- Here is one (simplified) way ...



Sample Locking System:

- (1) Don't trust transactions to request/release locks
- (2) Do not release locks until transaction commits/aborts:

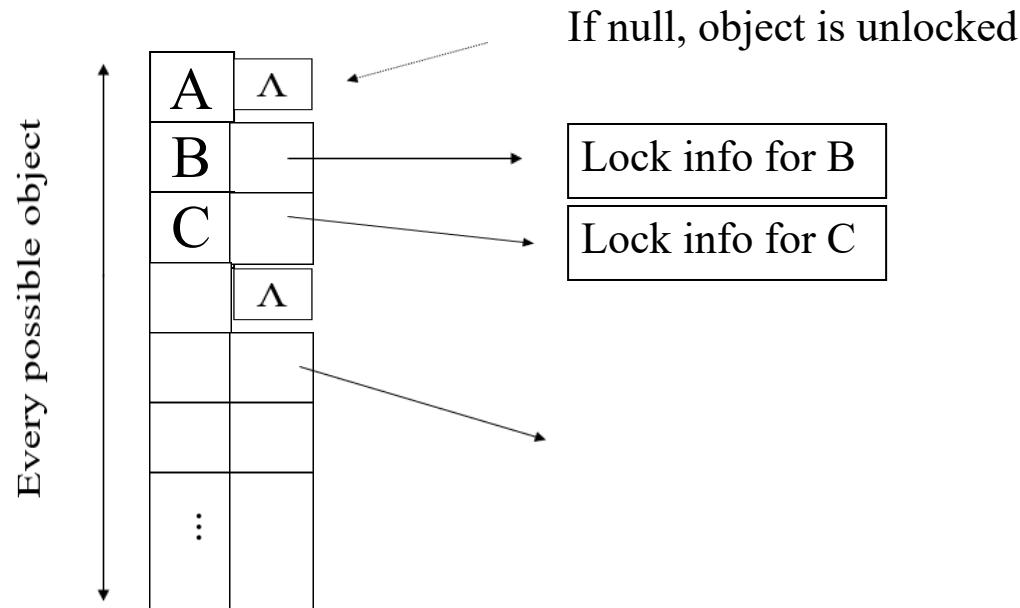




Database System - Nankai



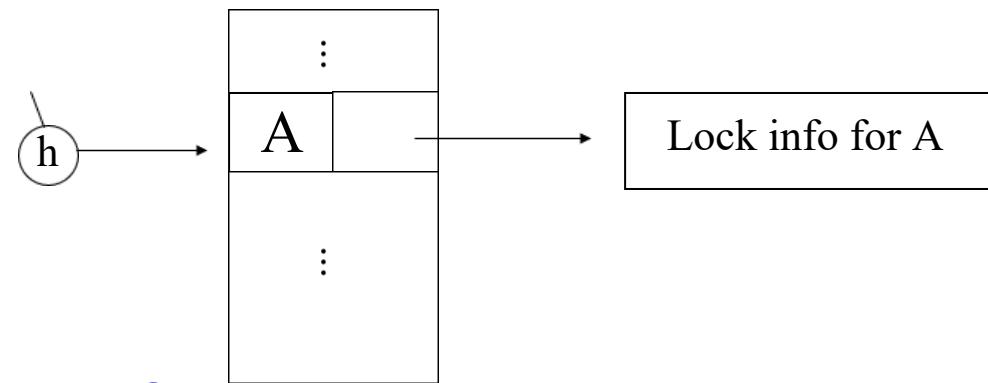
Lock table Conceptually





But use hash table:

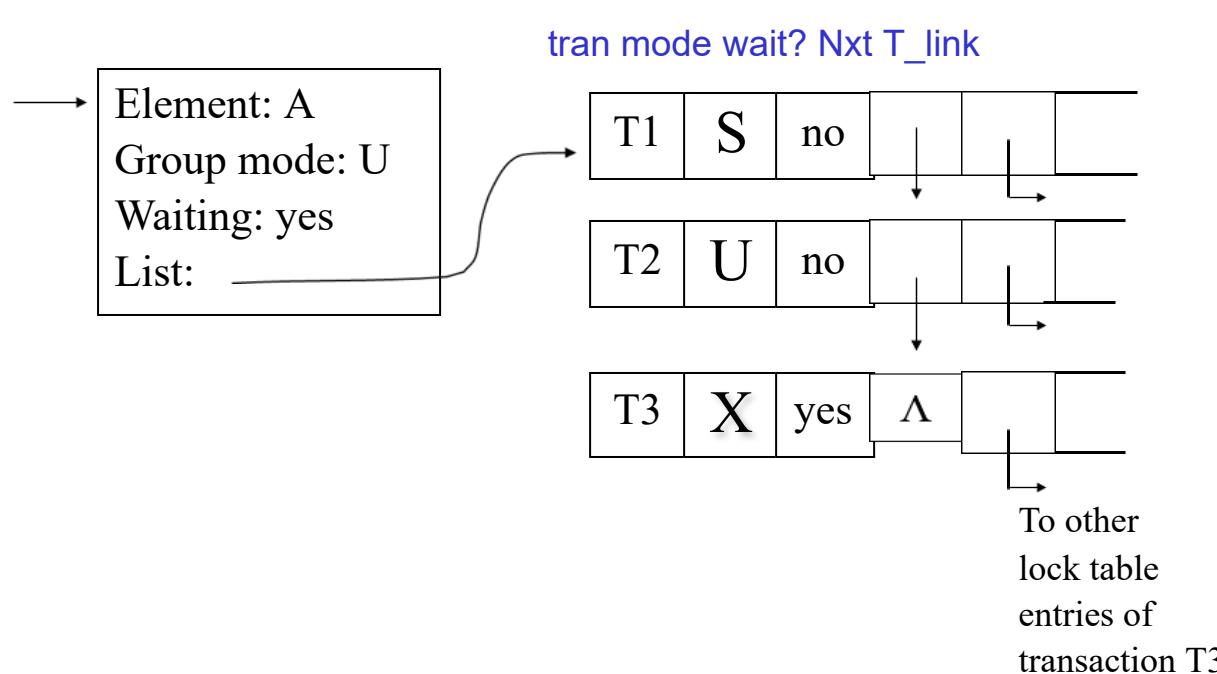
A



If object not found in hash table, it is unlocked



Lock info for A - example



多选题 1分



互动交流一

以下哪几个事务是Well-formed transactions并且满足两段锁协议？

- A $T = \text{ls}(A) \text{ lx}(B) r(A) w(B) u(B) \text{ lx}(C) w(C) u(C)$
- B $T = \text{ls}(A) r(A) \text{ lx}(B) w(B) u(B) \text{ lx}(C) w(C) u(C) u(A)$
- C $T = \text{ls}(A) r(A) u(A) \text{ ls}(B) w(B) u(B) \text{ lx}(C) w(C) u(C)$
- D $T = \text{ls}(A) \text{ lx}(B) r(A) w(B) \text{ lx}(C) u(B) w(C) u(C) u(A)$

提交

se System - Nankai



互动交流二 — 不定项选择题

以下哪几个调度满足合法调度(Legal scheduler)的规则?

- A $S = ls_1(A) r_1(A) ls_1(B) r_1(B) u_1(A) ls_2(B) r_2(B) u_2(B) ls_3(B) r_3(B) u_1(B)$
- B $S = ls_1(A) r_1(A) lx_1(B) w_1(B) ls_2(B) u_1(B) u_1(A) r_2(B) u_2(B) ls_3(B) r_3(B)$
- C $S = ls_1(A) r_1(A) ls_1(B) r_1(B) u_1(A) lx_2(B) w_2(B) u_2(B) ls_3(B) r_3(B) u_1()$

提交

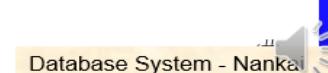
Database System - Nankai



Week16_Course

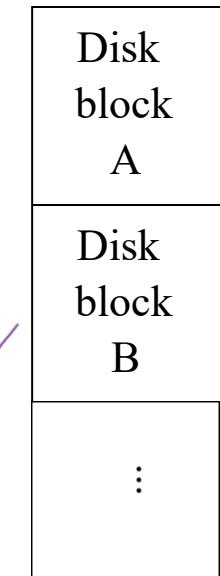
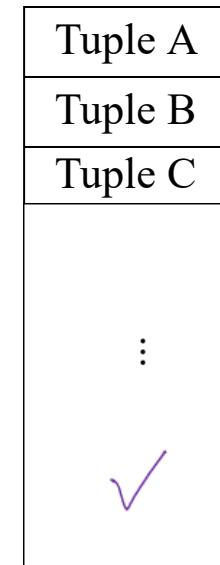
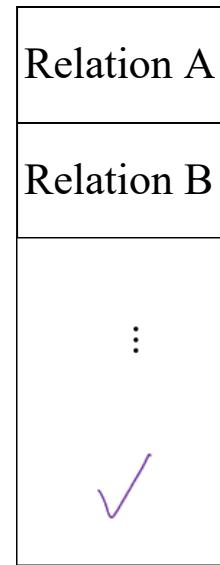
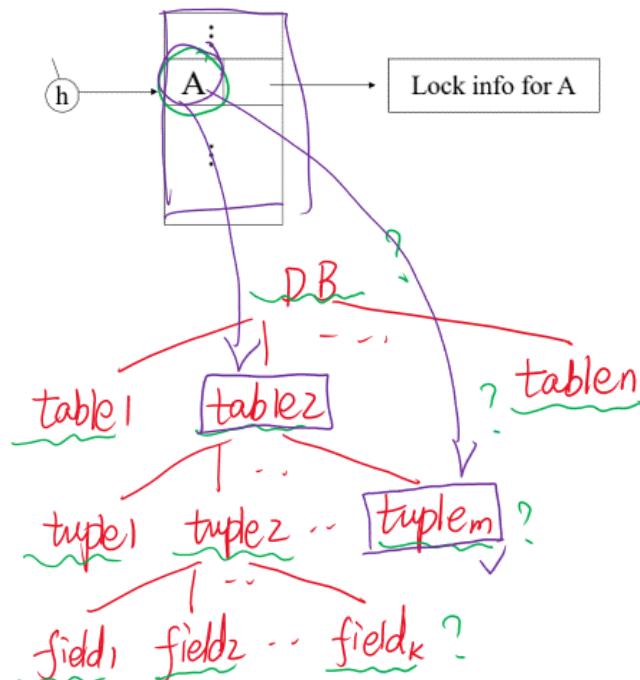
Database System Concurrency Control part4_2

Warning locks Protocol





What are the objects we lock?



?

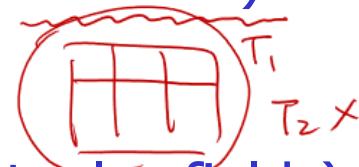
Database System - Nankai



- Locking works in any case, but should we choose small or large objects?

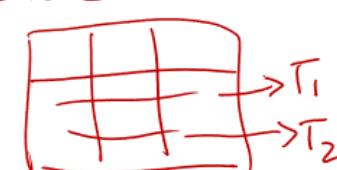
- If we lock large objects (e.g., Relations)

- Need few locks
 - Get low concurrency



- If we lock small objects (e.g., tuples, fields)

- Need more locks (\Rightarrow overhead higher)
 - Get more concurrency

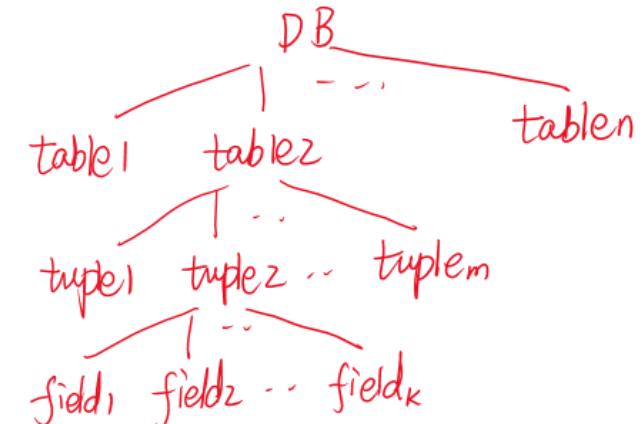


Database System - Nankai



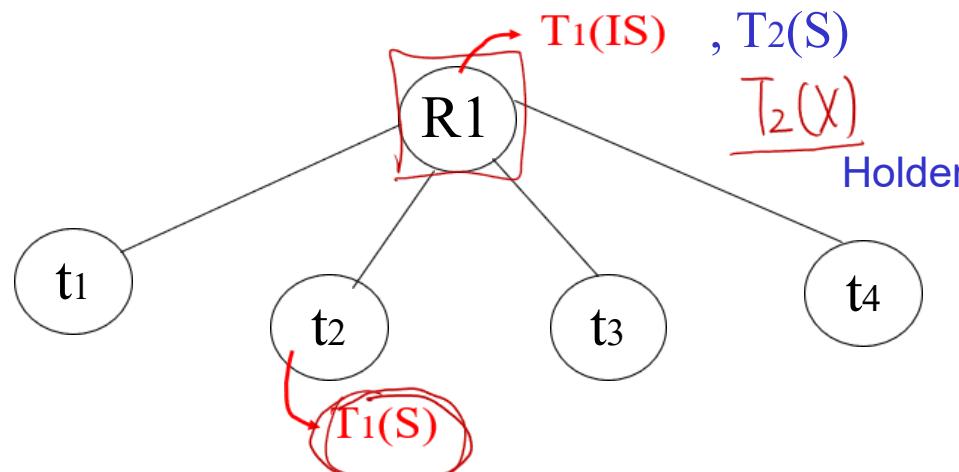
Warning Protocol

- Hierarchically nesting elements (e.g. relation/block/tuple) can be locked with **intention locks** IS and IX
- Idea
 - start locking at the root (relation) level
 - to place an S or X lock on a subelement, first place a corresponding intention lock IS or IX the element itself
 - Warns others: "I'll be reading/writing some subelement of this element"





Example (T_1 reads t_2 , T_2 ^{writes} R_1)



Comp Requestor

IS IX S X

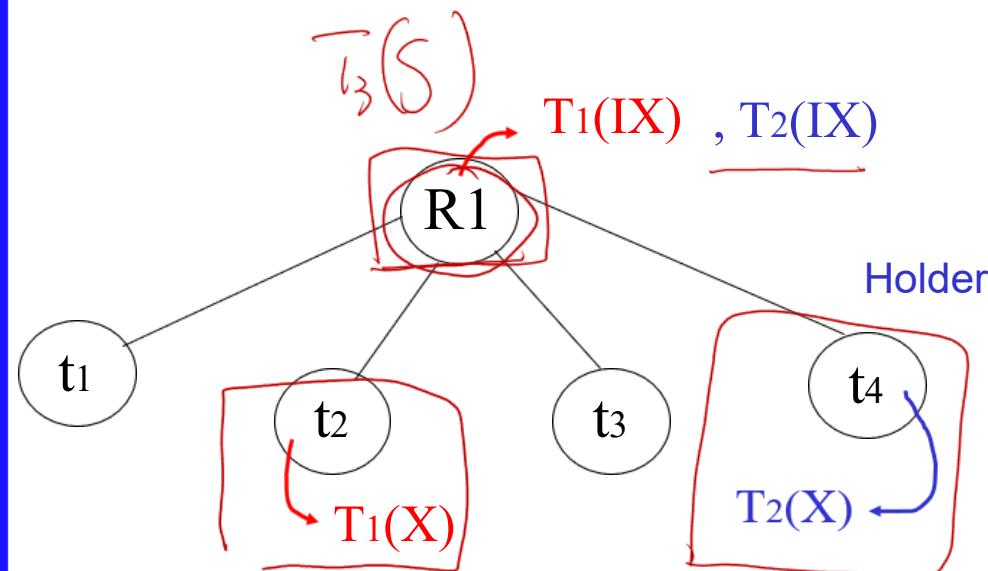
IS	T	T	T	F
S				
X				

As before

Database System - Nankai



Example (T_1 writes t_2 , T_2 writes to t_4)



Comp Requestor

	IS	IX	S	X
IS	T	T	T	F
IX	T	T	F	F
S	T	F		
X	F	F		

As before

Database System - Nankai



Compatibility of multiple granularity locks

Comp Requestor

IS IX S X

Holder IX

	IS	IX	S	X
IS	T	T	T	F
IX	T	T	F	F
S	T	F	T	F
X	F	F	F	F



SQL Server 完整的锁兼容性矩阵

	NL	SCH-S	SCH-M	S	U	X	IS	IU	IX	SIU	SIX	UIX	BU	RS-S	RS-U	RI-N	RI-S	RI-U	RI-X	RX-S	RX-U	RX-X
NL	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N
SCH-S	N	N	C	N	N	N	N	N	N	N	N	N	I	I	I	I	I	I	I	I	I	I
SCH-M	N	C	C	C	C	C	C	C	C	C	C	C	I	I	I	I	I	I	I	I	I	I
S	N	N	C	N	N	C	N	N	C	N	C	C	N	N	N	N	N	C	N	N	C	C
U	N	N	C	N	C	C	N	C	C	C	C	C	C	N	C	N	N	C	C	N	C	C
X	N	N	C	C	C	C	C	C	C	C	C	C	C	C	N	C	C	C	C	C	C	C
IS	N	N	C	N	N	C	N	N	N	N	N	C	I	I	I	I	I	I	I	I	I	I
IU	N	N	C	N	C	C	N	N	N	N	C	C	I	I	I	I	I	I	I	I	I	I
IX	N	N	C	C	C	C	N	N	N	C	C	C	I	I	I	I	I	I	I	I	I	I
SIU	N	N	C	N	C	C	N	N	C	N	C	C	I	I	I	I	I	I	I	I	I	I
SIX	N	N	C	C	C	C	N	N	C	C	C	C	I	I	I	I	I	I	I	I	I	I
UIX	N	N	C	C	C	C	N	C	C	C	C	C	I	I	I	I	I	I	I	I	I	I
BU	N	N	C	C	C	C	C	C	C	C	C	N	I	I	I	I	I	I	I	I	I	I
RS-S	N	I	I	N	N	C	I	I	I	I	I	I	I	N	N	C	C	C	C	C	C	C
RS-U	N	I	I	N	C	C	I	I	I	I	I	I	I	N	C	C	C	C	C	C	C	C
RI-N	N	I	I	N	N	N	I	I	I	I	I	I	I	C	C	N	N	N	N	C	C	C
RI-S	N	I	I	N	N	C	I	I	I	I	I	I	I	C	C	N	N	N	C	C	C	C
RI-U	N	I	I	N	C	C	I	I	I	I	I	I	I	C	C	N	N	C	C	C	C	C
RI-X	N	I	I	C	C	C	I	I	I	I	I	I	I	C	C	N	C	C	C	C	C	C
RX-S	N	I	I	N	N	C	I	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C
RX-U	N	I	I	N	C	C	I	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C
RX-X	N	I	I	C	C	C	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C	C

图例

N	不冲突	SIU	共享意向更新
I	非法	SIX	共享意向排他
C	冲突	UIX	更新意向排他
NL	没有锁	BU	大量更新
SCH-S	架构稳定性锁	RS-S	共享范围-共享
SCH-M	架构修改锁	RS-U	共享范围-更新
S	共享	RI-N	插入范围-空
U	更新	RI-S	插入范围-共享
X	排他	RI-U	插入范围-更新
IS	意向共享	RI-X	插入范围-排他
IU	意向更新	RX-S	持有范围-共享
IX	意向排他	RX-U	持有范围-更新
		RX-X	持有范围-Exclusive



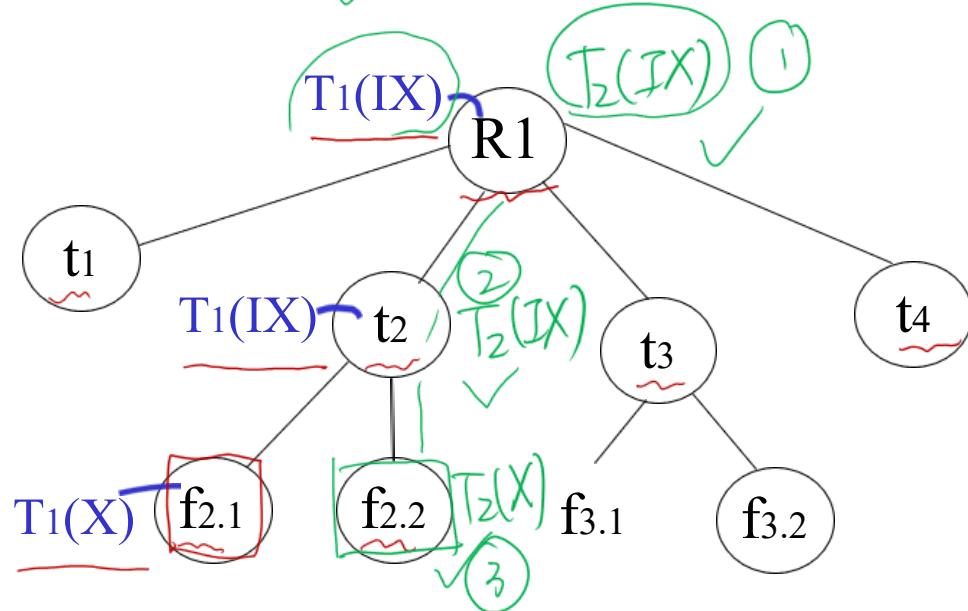
Rules

- (1) Follow multiple granularity comp function
- (2) Lock root of tree first, any mode
- (3) Node Q can be locked by Ti in S or IS only if
parent(Q) can be locked by Ti in IX or IS
- (4) Node Q can be locked by Ti in X,IX only if parent(Q)
locked by Ti in IX
- (5) Ti is two-phase
- (6) Ti can unlock node Q only if none of Q's children are
locked by Ti



Exercise:

- Can T₂ write element f_{2.2}? What locks will T₂ get?



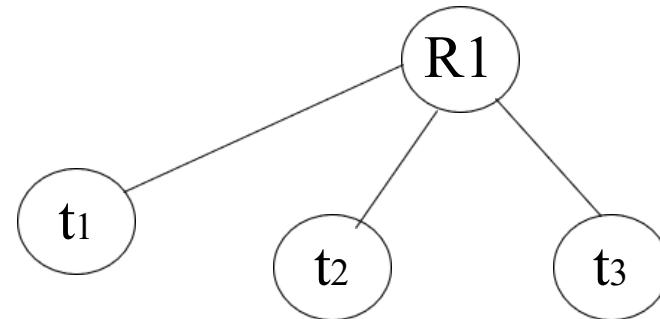
Database System - Nankai



Deletions and Phantom tuples

- Get an exclusive lock on A before deleting A
- Insertions more problematic:
 - possible to lock only existing elements
- **Phantom tuples:**
 - tuples that should have been locked, but did not exist when the locks were taken

- Use multiple granularity tree
- Before insert of node Q, lock parent(Q) in X mode



Database System - Nankai



Summary-more lock

- 了解锁的多种类型，重点掌握共享锁和排它锁，可以在一个事务中添加这些锁的动作，使其满足规则1和规则3
- 掌握带有读锁和写锁的事务如何生成合法的可串行化调度
- 了解DBMS系统是如何实现两段锁协议的
- 掌握意向锁的基本原理

多选题 1分



互动交流三 — 不定项选择题

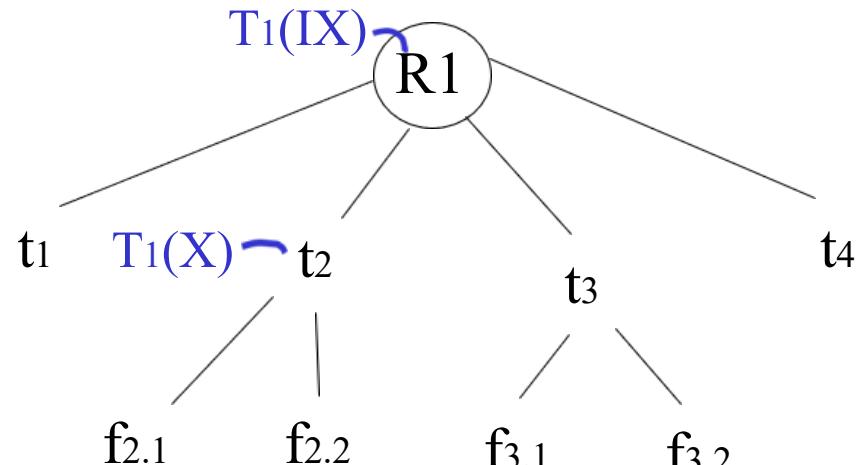
如果数据库元素的层次关系和持有锁的情况如图所示,

T₂事务要写f_{2.2}元素, 能否成功?

- A YES B NO

T₂事务要写f_{2.2}元素, 应该申请哪些锁?

- C R1的 IX锁
D t₂的 IX锁
E t₂的 X锁
F f_{2.2} 的X锁



提交

Database System - Nankai

多选题 1分



互动交流四 — 不定项选择题

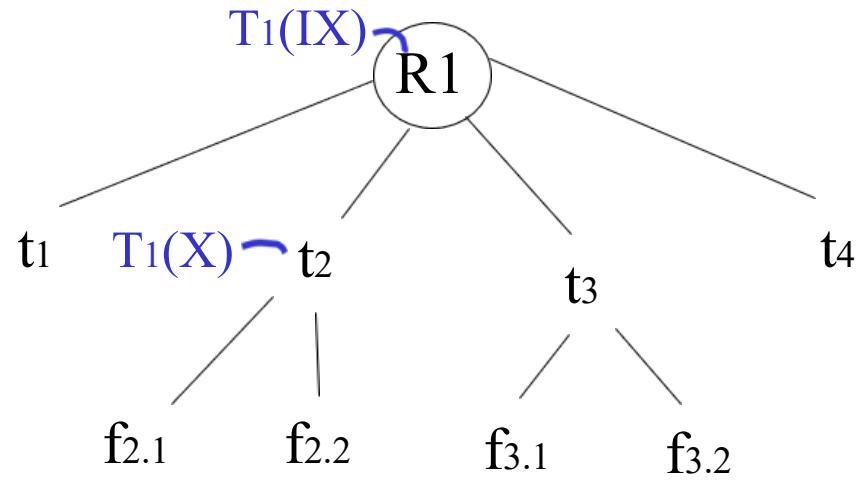
如果数据库元素的层次关系和持有锁的情况如图所示,

T₂事务要写f_{3.2}元素, 能否成功?

- A YES B NO

T₂事务要写f_{3.2}元素, 应该申请哪些锁?

- C R1的 IX 锁
D t₃的 IX 锁
E f_{3.2} 的 IX 锁
F f_{3.2} 的 X 锁



提交

Database System - Nankai

多选题 1分



互动交流五 — 不定项选择题

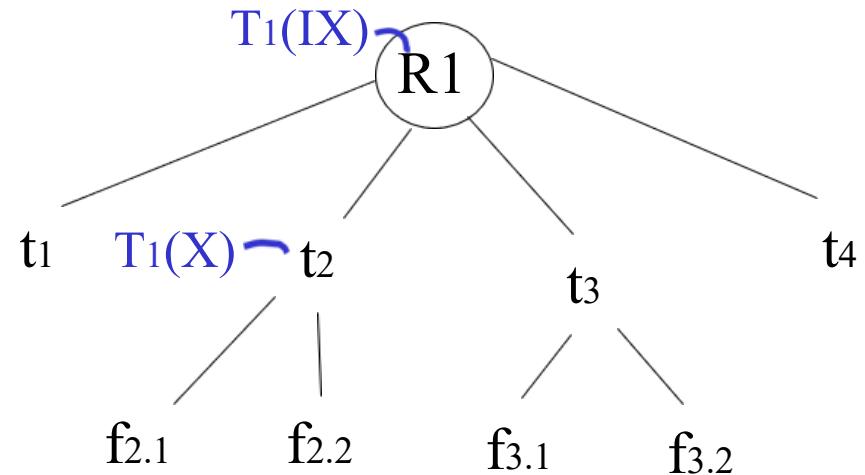
如果数据库元素的层次关系和持有锁的情况

如右图所示, T₂事务插入T₅元组, 能否成功?

- A YES B NO

T₂事务应该对R1申请什么锁?

- C R1的 IX 锁
D R1的 X 锁



提交

Database System - Nankai



Week16_Course

Database System Concurrency Control part5-1

Timestamp and Validation

Database System - Nankai



Concurrency Control

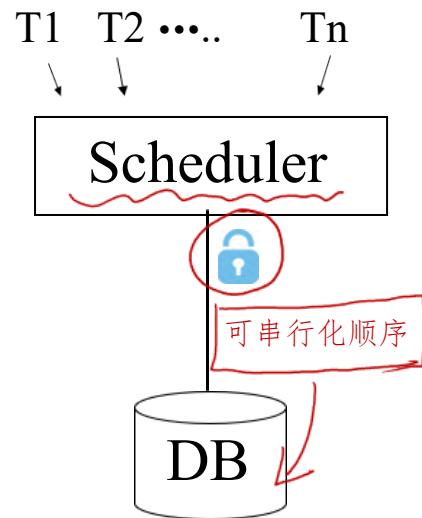
- Schedule
- Conflict Serializable Schedule
- Two phase locking
- Warning Protocol
-  Timestamp and Validation



防止出现不可串行化的主要技术

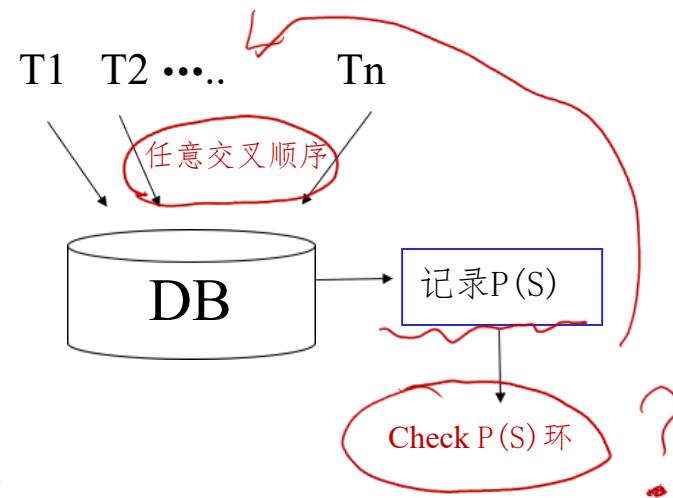
Pessimistic strategy: LOCK

Prevent occurrence of cycles in P(S)



Optimistic strategy

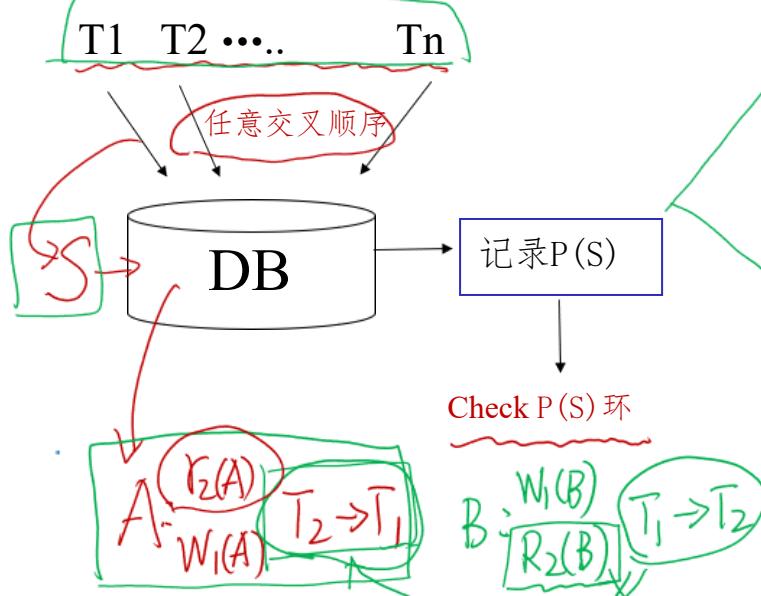
Run system, recording P(S); check P(S) for cycles,



Database System - Nankai



保证事务可串行化的其他方法



时间戳

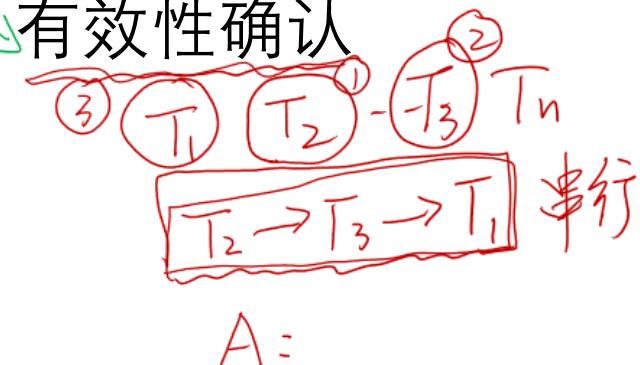
$$A: R_2(A): 3 \\ W_1(A): 5$$

$$T_1: 5 \quad T_2: 3 \quad T_3: 10$$

$T_2 \rightarrow T_1 \rightarrow T_3$

$$B: W_1(B): 5 \\ R_2(B): 3$$

有效性确认



A:

Database System - Nankai



保证事务可串行化的其他方法

- **时间戳**

- 每个事务分配一个“时间戳”
- 记录最后读和写每个数据库元素的事务时间戳
- 比较上述值以保证按照事务的时间戳排序的串行调度

- **有效性确认**

- 当事务将要提交时，检查事务的时间戳和数据库元素
- 按照有效性确认时间对事务进行排序的串行调度

中止并重启试图参与非可串行化行为的事务



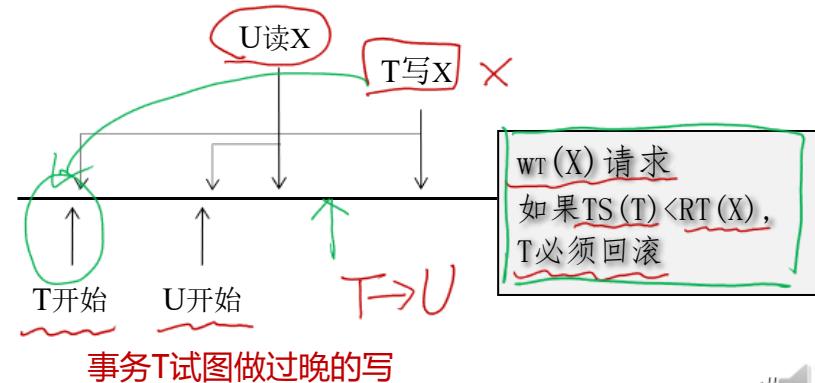
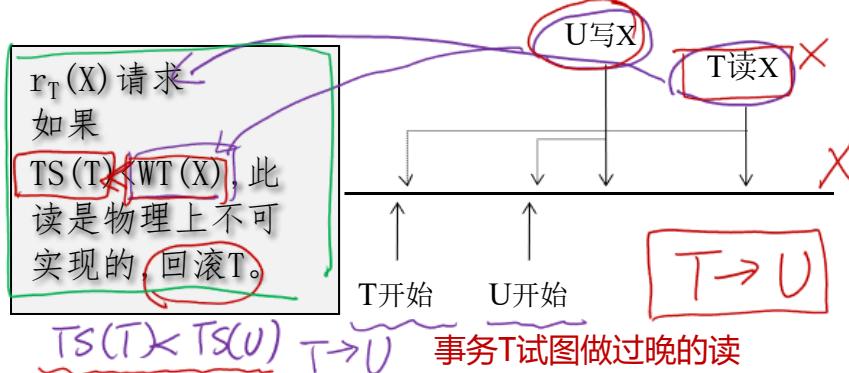


时间戳

$TS(T)$ T
 $TS(W)$ W
 $TS(K)$ K
 $RT - TS(W)$
 $X: WT - TS(K)$
 $C - T, F$

- 调度器赋给每个事务T一个唯一的数值 $TS(T)$
- 每个数据库元素X上有两个时间戳和一个附加位
 - $RT(X)$ 读X的事务中最高的时间戳
 - $WT(X)$ 写X的事务中最高的时间戳
 - $C(X)$ X的提交位, 该位为真当且仅当最近写X的事务已经提交

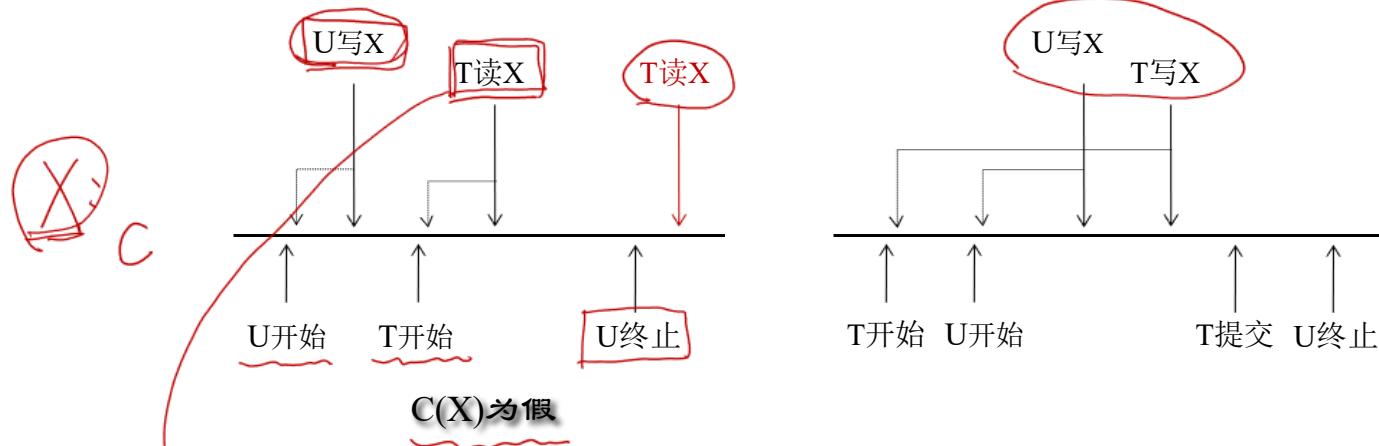
物理上不可实现的行为:



Database System - Nankai



脏数据问题



rt(X) 请求

- a. 如果 $TS(T) \geq WT(X)$, 此读是物理上可实现的
1. 如果 $C(X)$ 为真, 同意请求。如果 $TS(T) > RT(X)$, 置 $RT(X) := TS(T)$
 2. 如果 $C(X)$ 为假, 推迟 T 直到 $C(X)$ 为真或写 X 的事务终止

Database System - Nankai



基于时间戳调度的规则

1. $r_T(X)$ 请求

a. 如果 $TS(T) \geq WT(X)$, 此读是物理上可实现的

1. 如果 $C(X)$ 为真, 同意请求。如果 $TS(T) > RT(X)$, 置 $RT(X) := TS(T)$
2. 如果 $C(X)$ 为假, 推迟 T 直到 $C(X)$ 为真或写 X 的事务终止

a. 如果 $TS(T) < WT(X)$, 此读是物理上不可实现的。回滚 T。

1. $w_T(X)$ 请求

- a. 如果 $TS(T) \geq RT(X)$ 并且 $TS(T) \geq WT(X)$, 物理上可实现, 必须执行为 X 写入新值; 置 $WT(X) := TS(T)$ 并且置 $C(X) := false$
- b. 如果 $TS(T) \geq RT(X)$ 但是 $TS(T) < WT(X)$, 物理上可实现……
- c. 如果 $TS(T) < RT(X)$, T 必须回滚

1. 提交 T 请求…… $C(X) := true$

2. 终止 T 请求……

$r_T(X)$ 请求

- a. 如果 $TS(T) \geq WT(X)$, 此读是物理上可实现的
 1. 如果 $C(X)$ 为真, 同意请求。如果 $TS(T) > RT(X)$, 置 $RT(X) := TS(T)$
 2. 如果 $C(X)$ 为假, 推迟 T 直到 $C(X)$ 为真或写 X 的事务终止

$R_T(X)$ 请求

如果 $TS(T) < WT(X)$, 此读是物理上不可实现的, 回滚 T。

$w_T(X)$ 请求

如果 $TS(T) < RT(X)$, T 必须回滚



时间戳示例

T_2, T_3, T_1

<u>T_1</u>	<u>T_2</u>	<u>T_3</u>	<u>A</u>	<u>B</u>	<u>C</u>
200	150	175	R T=0 W T=0	R T=0 W T=0	R T=0 W T=0
r1(B)	r2(A)			R T=200	
w1(B)	r3(C)		R T=150		
w1(A)					
w2(C)			W T=200	W T=200	R T=175
终止			W T=200	✓	
	w3(A)	?			

(175) X

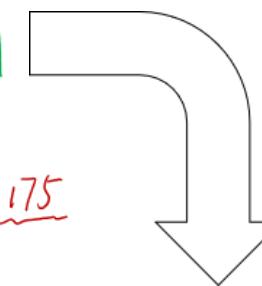
T_3

$T_2 = 150$



多版本时间戳

T1	T2	T3	T4	A
150	200	175	225	RT=0 WT=0
r1(A)				RT=150 WT=150
w1(A)				RT=200 WT=200
	r2(A)			
	w2(A)			
		r3(A)		
			r4(A)	RT=225



175

T1	T2	T3	T4	A ₀	A ₁₅₀	A ₂₀₀
150	200	175	225			
r1(A)				读		
w1(A)					创建	
	r2(A) 200				读	
	w2(A)					创建
		r3(A) 175		读		
			r4(A) 225			读

Database System - Nankai



时间戳与封锁

- 封锁在事务等待锁时会频繁地推迟事务
- 如果并发事务频繁读写公共元素，回滚会很频繁

商用系统中有一个折中方案，调度器将事务分为只读事务和读/写事务

- ✓ 只读事务使用多版本时间戳执行
- ✓ 读/写事务使用两阶段封锁执行

多选题 1分



互动交流——不定项选择题

以下哪些机制是DBMS采用的并发控制策略？

- A 采用两阶段锁的机制，控制并发事务的执行
- B 采用串行调度顺序，控制并发事务的执行
- C 按照事务的时间戳顺序，形成串行调度执行
- D 按照事务有效性确认顺序，形成串行调度执行

提交

se System - Nankai

多选题 1分



互动交流二

以下哪些策略，可能会造成数据库事务的频繁回滚？

- A 采用两阶段锁的机制，控制并发事务的执行
- B 采用串行调度顺序，控制并发事务的执行
- C 按照事务的时间戳顺序，形成串行调度执行
- D 按照事务有效性确认顺序，形成串行调度执行

提交

se System - Nankai

多选题 1分



互动交流三

以下哪种机制，会提高只读事务的并发度？

- A 两阶段封锁
- B 时间戳
- C 多版本时间戳
- D 任意机制都可以

提交

se System - Nankai